

Wine Quality Classification using Logistic Regression

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, accuracy_score, mean_
from math import sqrt

# Load dataset
df = pd.read_csv("winequality-red.csv", sep=";")

X = df.drop("quality", axis=1).values
y = (df["quality"] >= 6).astype(int).values.reshape(-1, 1) # Good

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Sigmoid function
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

# Logistic Regression
class LogisticRegression:
    def __init__(self, lr=0.01, epochs=2000):
        self.lr = lr
        self.epochs = epochs

    def fit(self, X, y):
        n_samples, n_features = X.shape
        self.weights = np.zeros((n_features, 1))
        self.bias = 0
        self.losses = []
```

```

for _ in range(self.epochs):
    # Linear model
    linear_output = np.dot(X, self.weights) + self.bias
    # Apply sigmoid
    y_pred = sigmoid(linear_output)

    # Compute loss (binary cross-entropy)
    loss = -(1/n_samples) * np.sum(
        y*np.log(y_pred+1e-9) + (1-y)*np.log(1-y_pred+1e-9)
    )
    self.losses.append(loss)

    # Gradients
    dw = (1/n_samples) * np.dot(X.T, (y_pred - y))
    db = (1/n_samples) * np.sum(y_pred - y)

    # Update weights
    self.weights -= self.lr * dw
    self.bias -= self.lr * db

def predict(self, X):
    linear_output = np.dot(X, self.weights) + self.bias
    y_pred = sigmoid(linear_output)
    return (y_pred >= 0.5).astype(int)

def predict_proba(self, X):
    linear_output = np.dot(X, self.weights) + self.bias
    return sigmoid(linear_output)

# Train model
model = LogisticRegression(lr=0.01, epochs=2000)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)

# Accuracy
accuracy = np.mean(y_pred == y_test)
print(f"Accuracy: {accuracy:.4f}")

# Calculate RMSE
rmse = sqrt(mean_squared_error(y_test, y_pred_proba))
print(f"RMSE: {rmse:.4f}")

# Plot loss curve
plt.figure(figsize=(7, 6))
plt.plot(model.losses)
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss Curve")
plt.grid(True)
plt.tight_layout()
plt.savefig('loss_curve.png')
plt.show()

# Calculate confusion matrix
cm = confusion_matrix(y_test, y_pred)

```

```


# Plot the confusion matrix
plt.figure(figsize=(7, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Bad Wine (0)', 'Good Wine (1)'],
            yticklabels=['Bad Wine (0)', 'Good Wine (1)'])
plt.title('Confusion Matrix - Wine Quality Classification')
plt.ylabel('Actual Quality')
plt.xlabel('Predicted Quality')
plt.tight_layout()
plt.show()

# plot of actual vs predicted probabilities
plt.figure(figsize=(7, 6))
plt.scatter(y_test, y_pred_proba, alpha=0.5)
plt.axhline(y=0.5, color='r', linestyle='--', lw=2, label='Decision')
plt.xlabel('Actual Quality (0=Bad, 1=Good)')
plt.ylabel('Predicted Probability of Good Wine')
plt.title('Actual vs Predicted Probabilities (RMSE: {:.3f})'.format(RMSE))
plt.yticks(np.arange(0, 1.1, 0.1))
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()


```

Accuracy: 0.7406

RMSE: 0.4134



Confusion Matrix - Wine Quality Classification



Actual vs Predicted Probabilities (RMSE: 0.413)

